

# Spatial Degradation of Hepatocyte Zonation in Human MASLD

An analysis report: building, validating and transferring a zonation ruler

2026-06-22

## 1 Background and question

The liver lobule is spatially organized along a porto-central axis: pericentral hepatocytes (near the central vein) and periportal hepatocytes (near the portal triad) run complementary metabolic programs, and most hepatocyte genes are *zonated*. **Paper 2** (Yakubovsky & Itzkovitz, *Nature* 2026) built a spatially resolved atlas of the *healthy* human liver and a quantitative zonation reference. **Paper 1** (Gribben, Vallier et al., *Nature* 2024, GSE202379) profiled ~69k hepatocytes by single-nucleus RNA-seq across 47 donors spanning MASLD stages (Healthy → NAFLD → NASH → Cirrhosis → End-stage), but is *spatially blind*. Paper 1 reports qualitatively that “hepatocytes lose their zonation”.

This project makes that quantitative. We use Paper 2’s healthy zonation as a *ruler* to assign each Paper 1 hepatocyte a pseudo-zonation coordinate, then test, at the level of the **donor** (never the cell, to avoid pseudoreplication):

- **H1**: does zonation collapse with disease stage?
- **H2**: which gene programs lose their zonal restriction (slope), and fastest?
- **H3**: are de-zonated cells also more plastic (hepatocyte↔cholangiocyte)?

## 2 Method, and the one rule that governs it

A ruler is any rule that maps a cell’s transcriptome to a porto-central coordinate. We compared many: Paper 2’s exact 20+20 *landmark* genes; curated anchor sets; data-driven top-*N* and “full” sets ranked from Paper 2’s zonation table; a *learned* dense PCA axis; sparse *lasso/elastic-net* axes; and a supervised whole-transcriptome classifier. Each yields  $\text{coord} = \text{mean}_Z(\text{PC}) - \text{mean}_Z(\text{PP})$  or a learned projection.

**The governing rule: ruler quality is judged ONLY on the healthy atlas** — an 8-marker sign gate (pericentral CYP2E1/CYP1A2/GLUL/PCK2 must correlate +; periportal ASS1/ALDOB/PCK1/HAL −), the healthy PC–PP anti-correlation (a real axis has the two arms mutually exclusive), and split-half reproducibility. The best healthy ruler is *frozen* and only *then* applied to the disease cohort. Disease H1/H2/H3 are the *test*; they are never used to choose a ruler. Choosing by disease strength would be circular (leakage).

## 2.1 Statistical tests — what each checks, and how

Every test follows the donor rule: collapse each donor’s cells to one number, then test on the  $\approx 47$  donor values (resampling/permuting donors, never cells). Tests are rank-based by default (ordinal stages, skewed scores, small  $n$ ); effect sizes are reported beside every  $p$ ; BH-FDR controls multiple testing.

| test                    | what it checks                          | how                                           |
|-------------------------|-----------------------------------------|-----------------------------------------------|
| Spearman $\rho$         | monotone trend of a metric vs stage     | rank correlation; donor-level                 |
| Jonckheere–Terpstra     | ordered-groups trend (dedicated)        | $J = \sum_{a < b} U_{ab}$ , permutation $p$   |
| Donor bootstrap CI      | how big / how sure the trend is         | resample donors $B=2000$ , 2.5/97.5 pct       |
| Label-shuffle perm.     | is the pipeline finding noise?          | shuffle stage labels, recompute               |
| Held-out gene split     | H2 without circularity                  | build coord on half genes, test other half    |
| Mann–Whitney $U$ (H2b)  | does a program weaken $>$ background?   | $U$ of program vs rest, rank-biserial, BH-    |
| Per-gene FDR (H2c)      | which individual genes lose slope       | Spearman per gene, BH over $\sim 30k$         |
| Within-donor corr (H3)  | dez $\leftrightarrow$ plast, stage-free | $\rho_d$ per donor; sign test & Wilcoxon vs 0 |
| Mann–Whitney AUC (H3)   | dez vs zonated plasticity               | within-donor AUC, aggregate                   |
| OLS + C(stage)+C(donor) | partial dez effect (descriptive)        | coefficient on dez (cell-level SE)            |

## 3 Results and analysis

### 3.1 Not every ruler is a ruler: a dilution curve

The 8-marker gate passes for almost everything (the markers are strong), but it is not sufficient. The decisive control is the *healthy PC–PP anti-correlation*. Small, high-SNR sets are genuinely bipolar (`paper2.landmark`:  $-0.45$ ; `expanded_curated`:  $-0.48$ ), but as the gene set grows the axis degrades and finally inverts: `paper2.top250`  $\approx +0.04$ , `paper2.full`  $\approx +0.55$ . Averaging  $\sim 2000$  weakly-zonated genes lets a shared technical/expression factor dominate, so the “full” coordinate is *not* a zonation axis at all. This is why a transcriptome-wide *unweighted* mean fails while a *weighted* or selected ruler succeeds.

### 3.2 Ruler ranking and selection (healthy metrics ONLY)

Every ruler is scored on **healthy** quality alone — a quality score of (split-half reproducibility)  $+ \max(0, -\text{healthy PC–PP anticorrelation})$  — and the table below is **ordered by it**. Disease-collapse strength is deliberately *excluded* from selection: using the test cohort to choose the instrument would be leakage.

**Eligibility is by leakage, not by publishability.** A ruler may compete for the frozen primary slot iff its axis was *not* fit on Paper-1 cells. This matters because two learned rulers (`unsupervised`, `unsupervised_combined`) are fit by PCA on Paper-1 healthy cells — so their healthy anticorrelation/split-half are computed on the very cells the axis was fit to (*in-sample*, and therefore optimistically inflated; that is exactly why they top the raw table). They remain valid H1 robustness checks — the Paper-1 *disease* cells are never used to fit them — so we keep and show them, marked *control* (*Paper1-fit*), but they do not compete. Every other ruler (published gene lists, which fit nothing; and the Paper-2-trained learned axes `unsupervised_p2/supervised/lasso/elasticnet`, external to Paper 1) is leakage-clean and eligible regardless of whether it is a curated list or a label-free axis.

| ruler                   | PC   | PP   | healthy anticorr | split-half | H1 spread $\rho$ | role                    |
|-------------------------|------|------|------------------|------------|------------------|-------------------------|
| unsupervised_combined   | 4769 | 3719 | -0.50            | 0.72       | -0.51            | control (Paper1-fit)    |
| unsupervised_p2         | 5177 | 3302 | -0.47            | 0.69       | -0.50            | PRIMARY (label-free)    |
| expanded_curated        | 26   | 23   | -0.48            | 0.67       | -0.43            | PRIMARY (interpretable) |
| selected_frozen         | 26   | 23   | -0.48            | 0.67       | -0.43            | primary_candidate       |
| unsupervised_top100     | 100  | 100  | -0.29            | 0.82       | -0.61            | primary_candidate       |
| unsupervised_top50      | 50   | 50   | -0.31            | 0.79       | -0.62            | primary_candidate       |
| unsupervised            | 4585 | 3903 | -0.31            | 0.78       | -0.56            | control (Paper1-fit)    |
| paper2_landmark         | 20   | 20   | -0.45            | 0.63       | -0.38            | primary_candidate       |
| unsupervised_top250     | 250  | 250  | -0.23            | 0.85       | -0.57            | primary_candidate       |
| paper2_top50            | 50   | 50   | -0.40            | 0.57       | -0.32            | primary_candidate       |
| paper2_top100           | 100  | 100  | -0.27            | 0.56       | -0.28            | primary_candidate       |
| unsupervised_full       | 4585 | 3903 | 0.65             | 0.83       | -0.42            | exploratory             |
| core_curated            | 13   | 8    | -0.28            | 0.54       | -0.54            | interpretability_only   |
| unsupervised_lasso      | 126  | 149  | -0.18            | 0.63       | -0.51            | robustness              |
| unsupervised_elasticnet | 325  | 281  | -0.07            | 0.69       | -0.49            | robustness              |
| paper2_top250           | 250  | 250  | 0.04             | 0.55       | -0.27            | exploratory             |
| paper2_full             | 1624 | 447  | 0.56             | 0.48       | -0.04            | exploratory             |
| full                    | 0    | 0    | 0.53             | 0.46       | 0.02             | exploratory             |
| supervised              | 4504 | 3935 | -0.01            | 0.45       | 0.12             | exploratory             |

### 3.3 Two co-primary rulers, and why we report both

Among the leakage-clean rulers we freeze **two co-primaries** of deliberately different construction, then transfer *both* to disease: an **interpretable** published anchor (**expanded\_curated**) and a **label-free** learned axis (**unsupervised\_p2**, trained entirely on the external Paper 2 healthy atlas). On healthy quality they are statistically indistinguishable (healthy-quality scores 1.15 vs 1.16, a gap of 0.012 — within the 0.02 tie band), so picking one over the other would be arbitrary — and the label-free axis is in fact the *highest*-scoring eligible ruler, so the result is not an artefact of insisting on a hand-curated gene list. The headline rests on the two *agreeing*: a curated signature and an unsupervised axis built with no marker genes both reconstruct the same coordinate and both show the same collapse. We also note the selected rulers are *not* the steepest-collapsing ones (the tiny curated **core\_curated** set collapses harder) — choosing by disease strength would be cherry-picking; selection used healthy metrics only, and disease is the independent confirmation.

### 3.4 H1 — zonation collapses (and it is robust)

On every *valid* ruler, per-donor coordinate spread falls and the PC–PP anti-correlation rises toward zero across stages. The signal is concentrated and reproducible: landmark spread  $\rho = -0.380$  (perm  $p = 0.0085$ ); expanded  $\rho = -0.434$ ; the label-free dense PCA axis  $\rho = -0.504$  (perm  $p = 0.0005$ ). The broken “full” ruler is null ( $\rho = -0.042$ ), exactly as its healthy diagnostics predicted. The agreement of the published-signature ruler and a *label-free* ruler is the key point: the collapse is a property of the data, not of any one gene list.

**Early disease is the noisy part (stated honestly).** The decline is *not* strictly monotonic from Healthy: at NAFLD the coordinate is often as zoned as (or slightly more than) Healthy, and the Healthy→NAFLD *direction flips across rulers* (up for landmark/expanded, down for the dense PCA axis) — exactly the instability expected where the signal is weakest and  $n$  smallest (Healthy

$n = 4$ , NAFLD  $n = 7$ ). All rulers nonetheless *agree* on the strong NASH→end-stage decline. So the defensible claim is an ordered collapse *across the disease course, driven by NASH→end-stage*, with early steatosis leaving zonation largely intact — consistent with steatosis beginning pericentrally before inflammation and fibrosis dissolve the axis.

### 3.5 Leakage control — the unsupervised axis is clean

A fair concern: was the learned axis trained on disease cells? No. It was learned on healthy cells only and frozen. To remove all doubt we also trained the axis *entirely on the external Paper 2 healthy atlas* (zero Paper 1 cells in fitting) and transferred it: the collapse survived ( $\rho = -0.50$ ,  $p = 3 \times 10^{-4}$ , 8/8 markers). And the label-free PCA axis correlates  $\rho = +0.57$  with the landmark-signature coordinate ( $\rho = +0.39$  with its PC arm,  $-0.45$  with its PP arm) — so the axis is not a signature-formula artefact, and the result is not a variance-maximization artefact.

### 3.6 H2 — which programs lose zonation, and fastest

**H2a (held-out slope-loss).** Building the coordinate from one random half of the ruler genes and testing the other half (repeated  $K=20\times$ ), most held-out genes lose their donor-level zonal slope with disease (e.g. 96/100 weaken for the unsupervised ruler) — a non-circular confirmation that the genes defining zonation flatten with disease.

**H2c (transcriptome-wide, per gene).** Taking it further: assign every cell the frozen ruler coordinate and test *all* 29969 genes’ donor-level zonal slope. 67% lean toward weakening (a real directional bias — pure noise would give 50%), *but only 8 survive BH-FDR* at  $q < 0.05$ . Why so few? Each gene’s slope is estimated per donor from a handful of zone bins, and with only  $n=47$  donors the per-gene test is weak; correcting across  $\sim 30k$  genes, almost no single gene is individually a stand-out. So the collapse is **broad and diffuse** — most genes lose a *little* zonation rather than a few master genes losing a lot.

**H2b (program-level — where the power and the structure are).** Diffuse does *not* mean uniform. Grouping genes into programs and aggregating recovers the power that per-gene FDR throws away: for each program we test, by Mann–Whitney  $U$ , whether its genes weaken *more than the genome background* (background median  $\rho = -0.06$ , 67% weakening), with a rank-biserial effect and BH across programs. The programs weaken at *significantly different rates*:


| program                  | $n$   | median $\rho$ | MWU $q$ vs background |
|--------------------------|-------|---------------|-----------------------|
| wnt_pericentral_identity | 11    | -0.40         | 8.5e-06               |
| urea_nitrogen            | 14    | -0.21         | 3.1e-03               |
| cyp_xenobiotic           | 30    | -0.20         | 5.4e-05               |
| lipid_metabolism         | 20    | -0.18         | 0.01                  |
| bile_acid_transport      | 11    | -0.12         | 0.55                  |
| cholangiocyte_plasticity | 13    | -0.11         | 0.60                  |
| other                    | 29845 | -0.06         | —                     |
| acute_phase_secreted     | 25    | +0.04         | 1.00                  |

Pericentral Wnt/identity genes de-zonate first and hardest (median  $\rho = -0.40$ , MWU  $q \approx 8e - 06$ , rank-biserial  $-0.82$ ), followed by urea/nitrogen, xenobiotic CYP and lipid metabolism (all  $q < 0.05$  more negative than background); acute-phase/secreted and bile-acid programs are *spared* (not different from, or less negative than, background). This is the resolution of the apparent paradox: per-gene FDR (H2c) flags few individuals because each is noisy, yet the *between-program* contrast

(H2b) is highly significant because averaging within a program and comparing against the rest is far better powered. Biologically coherent — loss of pericentral Wnt identity leads, while acute-phase genes change *globally* (not zonally) so their zonal slope is untouched. (H2 = loss of spatial restriction, distinct from classic level-DE; cross-tabulated per gene in `h2_slope_loss_vs_DE.csv`.)

### 3.7 H3 — de-zonation couples weakly to plasticity

Within donors, de-zonation correlates positively with plasticity-marker expression (KRT7/KRT19/SOX9/SOX4/KRT18) in  $\sim 70\%$  of donors (landmark sign-test  $p = 0.008$ ), but the effect is small (mean  $\rho \approx 0.02$ ). This is the weakest of the three hypotheses: directionally consistent with Paper 1’s transdifferentiation theme, but not a strong donor-level effect here.



## 4 Discussion

Three independent constructions — a published landmark signature, a label-free PCA axis (trained even on a fully external healthy atlas), and regularized sparse axes — converge on the same conclusion: hepatocyte zonation progressively collapses across MASLD, the pericentral Wnt-identity program leads the loss, and de-zonation is mildly associated with epithelial plasticity. Because ruler selection was confined to healthy-atlas criteria, the disease findings are genuine transfer results rather than fitted outcomes.

## 5 Limitations

- (i) Donor numbers per stage are imbalanced (4/7/27/4/5), so cirrhosis/end-stage estimates are thin; all inference is donor-level with bootstrap CIs and permutation nulls to respect this.
- (ii)

Paper 2 per-cell zone labels are an  $\eta$ -over-landmark reconstruction (Paper 2’s own method) — spatially grounded but not a physical coordinate; classifier entropy is therefore auxiliary. (iii) The “full” unweighted set is a cautionary negative, not a result. (iv) H2b program groupings are curated and the smaller rulers contain few genes per program, so the gene-rich **paper2\_full** table is used for the program comparison.

## 6 Conclusions

Hepatocyte zonation collapses with MASLD stage at the donor level (H1), driven first by the pericentral Wnt-identity program with nitrogen and xenobiotic metabolism following (H2), and is weakly coupled to epithelial plasticity (H3). The result is robust to how the ruler is built and to training the ruler on a fully external healthy atlas.

Companion artefacts: per-set dossier **Zonation\_Ruler\_Report.pdf**; machine tables **signature\_battery\_summary.csv**, **selected\_set\_decision.txt**, per-set **results/tables/<set>/**.